

COMP 364 / 464

High Performance Computing

OpenMP:

Multi-threaded and task parallelism

Computing Terminology

Terms	Definition
• NUMA	• Non Uniform Memory Access. In SMP systems with multiple CPUs, access time to different parts of memory may vary
• Affinity	• Propensity to maintain a process or thread on a hardware execution unit
• SMP	• Symmetric Multi-Process(ing/or). Single OS system with shared memory
• OpenMP	
– Directive	• Comment statement (F90) or #pragma (C/C++) that specifies parallel operations and control
– Construct	• The lexical extent that a directive controls
– Region	• All code controlled by a directive– lexical extent + content of called routines
• Runtime	• Code or a library within an executable that interacts with the operating system and can control code execution

What is OpenMP (*Open Multi-Processing*)

- De-facto standard for high-performance parallel programming on *Symmetric Multi-Processor* (SMP) Systems
- It is an API (*Application Program Interface*) for designing and executing parallel Fortran, C, and C++ programs
 - Based on threads, but
 - Higher level than POSIX threads (*Pthreads*)
<http://www.llnl.gov/computing/tutorials/pthreads/#Abstract>
- Implemented by:
 - Compiler Directives (or #pragma's)
 - Runtime Library (interface to OS and Program Environment)
 - Environment Variables
- Compiler option required to interpret/activate directives
- <http://www.openmp.org/> has tutorials and description
- Directed by OpenMP ARB (*Architecture Review Board*)

OpenMP -- Overview

- Standard developed in the late 1990s
- The “language” is easily grasp. You can start simple and expand.
- Lightweight from system perspective
- Very portable – GNU GCC and vendor (Intel) compilers.
- OMP compilers generate the parallel code but developers must define the parallelism and communication mechanisms.
 - Time spent finding parallelism can be the most difficult part. The parallelism may be hidden.
 - Writing parallel OpenMP code examples is relatively easy.
 - Developing parallel algorithms and/or parallelizing existing serial code is much harder.
 - Expert level requires awareness of scoping and synchronization.

OpenMP Executable Runs on an SMP*

- Shared Memory Systems:
 - One Operating System
 - Instantiation of ONE process
 - Threads are forked (created) from within your program
 - Multiple threads on multiple cores

* SMP = Symmetric Multi-Processor: The execution of the operating system has equal access to any of the “processors”

OpenMP History

- Primary OpenMP participants

AMD, Cray, Fujitsu, HP, IBM, Intel, NEC, PGI, Oracle, MS, TI, CAPS, NVIDIA
ANL, LLNL, cOMPunity, EPCC, LANL, NASA, ORNL, RWTH, TACC

- | | |
|--------------------------------------|------------------------|
| • OpenMP Fortran API, Version 1.0 | Published October 1997 |
| • OpenMP C API, Version 1.0 | Published October 1998 |
| • OpenMP 2.0 API for Fortran | Published 2000 |
| • OpenMP 2.0 API for C/C++ | Published 2002 |
| • OpenMP 2.5 API for C/C++ & F90 | Published 2005 |
| • OpenMP 3.0 Tasks | Published May 2008 |
| • OpenMP 3.1 | Published July 2011 |
| • OpenMP 4.0 Affinity, Offload, SIMD | Published July 2013 |
| • OpenMP 4.5 | Published Nov 2015 |

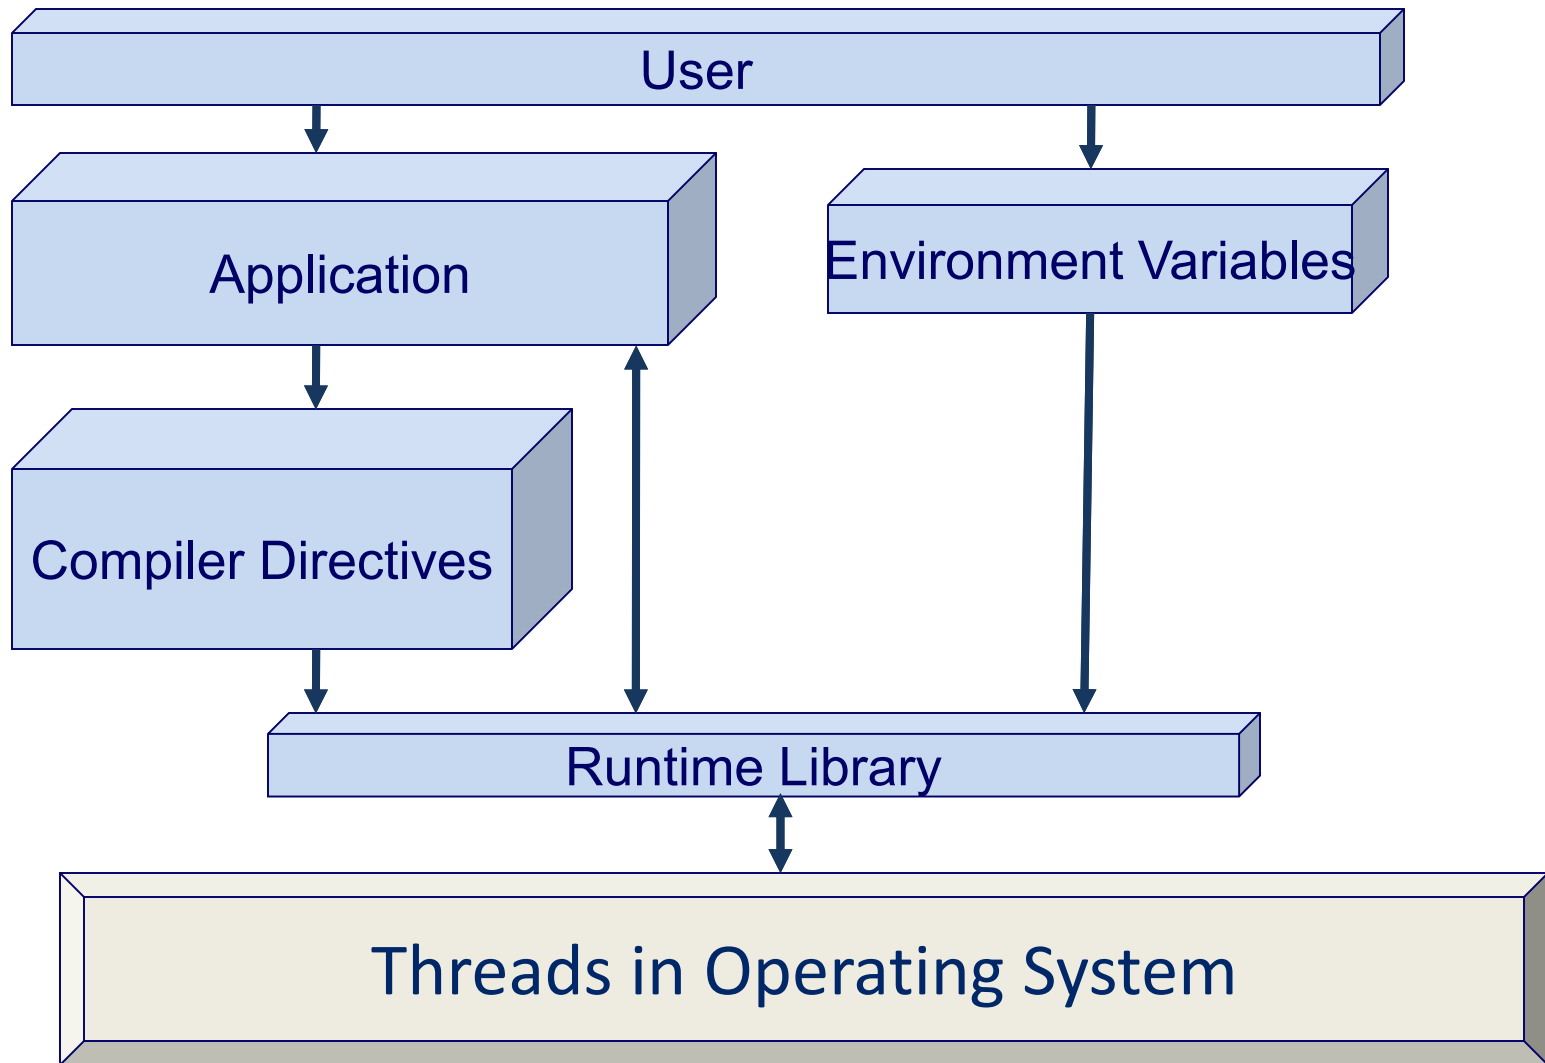
Advantages/Disadvantages of OpenMP

- Pros
 - Shared Memory Parallelism is easier to learn than distributed MP
 - Coarse-grained or fine-grained parallelism
 - Parallelization can be incremental
 - Widely available, portable
 - Converting serial code to OpenMP parallel can be easier than converting to multi-process, distributed parallel
 - SMP hardware is ubiquitous now: Supercomputers **and** your desktop/laptop
- Cons
 - Scalability limited by memory architecture:
 - Non-uniform memory access (NUMA) can be devastating.
 - Thread synchronization dependent upon fast access to global variables – in main memory.
 - Available on SMP systems “only”
 - Beware: “Upgrading” large serial code may be time-consuming ... but it always will be!
 - Not directly part of the language.

OpenMP Parallel Directive

- Supports parallelism by Directives (Fortran) and Pragmas (C/C++)
- Unlike others that require base language changes and constructs (Pthreads)
- Unlike MPI which supports parallelism through communications library between processes.

OpenMP Architecture



OpenMP Syntax

- OpenMP Directives: **sentinel**, **construct** and **clauses**

```
C      #pragma omp construct [clause [[,]clause]...]
```

```
F90    !$omp          construct [clause [[,]clause]...]
```

- Example

```
C      #pragma omp parallel num_threads(4)
```

```
F90    !$omp          parallel num_threads(4)
```

- Function prototypes and types are in the file:

```
C      #include <omp.h>
```

```
F90    use omp_lib
```

- Most OpenMP constructs apply to a “structured block”, that is, a block of one or more statements with one point of entry at the top and one point of exit at the bottom

OpenMP Constructs

OpenMP Language “extensions”

Parallel Control Structures

- governs flow of control in the program

parallel directive

Parallel Control worksharing

- distributes work among threads

do
for
sections
single
construct

Control Single Task

- assigns work to a thread

task
taskgroup
construct

Data Environment

- specifies variables as shared or private

shared
private
reduction
clause

Synchronization

- coordinates thread execution

critical
atomic
barrier
taskwait
directive

Runtime Environment

- runtime functions
- environment variables

omp_set_num_threads()
omp_get_thread_num()
OMP_NUM_THREADS
OMP_SCHEDULE

- **scheduling**
static, dynamic, guided

Directive *Sentinels*: “!\$omp” and “#pragma omp” not shown.

OpenMP Directives

- OpenMP directives begin with special comments/pragmas that a OpenMP-aware compiler can interpret. Directive sentinels are:

F90 **! \$OMP**

C/C++ **# pragma omp**

Syntax: *sentinel parallel clauses* *uses defaults when clauses not present*

Fortran

```
!$OMP parallel
...
!$OMP end parallel
```

C/C++

```
# pragma omp parallel
{ ... }
```

- Fortran parallel regions are enclosed by enclosing directives.
- C/C++ parallel regions are enclosed by curly brackets.

Parallel Region & Thread Number

```
1 #include <omp.h>
2 ...
3 int ierr = 0;
4 #pragma omp parallel
5 {
6     int numThreads = omp_get_num_threads();
7     int threadId = omp_get_thread_num();
8     ierr = do_work( threadId, numThreads );
9 }
```

Every thread can inquire the total number of threads (**numThreads** in line 6) and get a unique value for the thread number [0,K-1)

Parallel Region & Thread Number

```
1  #if defined(_OPENMP)
2  # include <omp.h>
3  #endif
4  ...
5  int ierr = 0;
6  #pragma omp parallel
7  {
8      int numThreads = 1;
9      #if defined(_OPENMP)
10         numThreads = omp_get_num_threads();
11     #endif
12     ierr = work(numThreads);
13 }
```

OpenMP code can be disabled and ignored by the compiler. All sentinels and directives ignored when OMP isn't enabled by compiler. Pre-processor (macro) flag useful to remove OMP API function calls and header. `_OPENMP` automatically defined when OMP enabled.

Parallel Region & Worksharing

Use OpenMP directives to specify Parallel Region, *worksharing* constructs, and mutual exclusion

parallel
end parallel

Use parallel ... end parallel for F90
Use parallel {...} for C

parallel for
parallel sections

Code block
for
sections
single
master
critical
atomic

Each Thread Executes
Worksharing: splice loop
Worksharing: splice calls
Only one thread (first there)
Only master thread (0)
Execute one thread at a time
Update one thread at a time

A single worksharing construct
(e.g. **for**) may be combined
on a parallel directive line.

Parallel Region

```
...  
#pragma omp parallel  
{  
    #pragma omp for  
    for ( i = 0; i < n; i++){  
        work(i);  
    } // <- implicit barrier  
} // <- implicit barrier
```

- In above example the **for** loop iterations are split among the threads via the **for** worksharing construct.
- Which iteration(s) are executed by which threads? Controlled by **schedule** construct. Default is to split the iterations into $K=\text{\#threads}$ chunks, each n/K long.

OpenMP Combined Directives

- Combined directives
 - **parallel for** and **parallel sections**
 - Same as parallel region + one **for** worksharing construct

```
#pragma omp parallel for  
for (i= 0; i< 100; i++){  
    a[i] = b[i];  
} // <- implicit barrier
```

- Fixed (known) trip count and increment required
- no break
- limited C++ throw
- continue ok (since it doesn't impact other iterations)
- Basic rule of thumb: if it can be done with SIMD, it can be threaded with OpenMP-for

Parallel Region

worksharing (WS) constructs: **for**, **sections**, and **single**

- WS Threads execution their “share” of statements in a PARALLEL region.
- **for** worksharing may require run-time work distribution and scheduling

```
1 #pragma omp parallel for
2 for (i= 0; i< n; i++){
3     a[i] = b[i] + c[i];
4 } // <- implicit barrier
```

Line 1: Team of threads formed (parallel region).

Line 2-4: Loop iterations are split among threads.

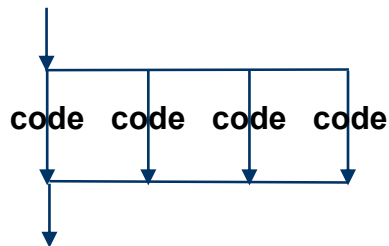
Implied barrier at “}”.

Each loop iteration must be independent of other iterations!

Replicated vs. Workshare Constructs

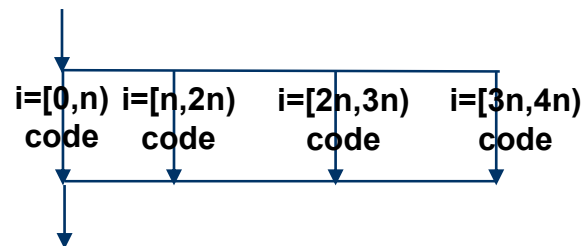
- Replicated: Work (code) blocks are executed by all threads.
- Worksharing: Work is divided among threads ... no overlap.

```
#pragma omp parallel  
{code}
```



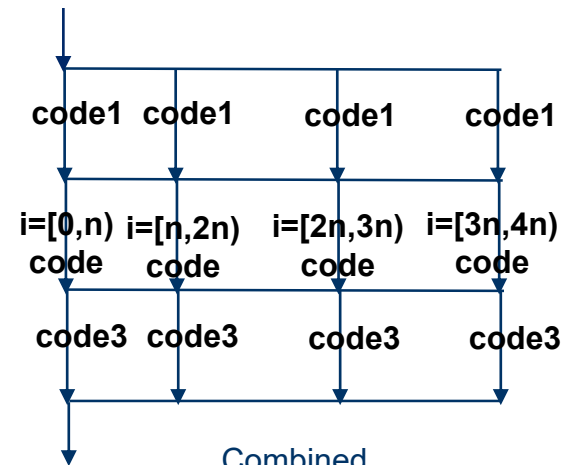
Replicated

```
#pragma omp parallel for  
for (i=0; i<4*n; i++)  
{code}
```



worksharing

```
#pragma omp parallel  
{  
    code1();  
    #pragma omp for  
    for (i=0; i<4*n; i++)  
        code();  
    code3();  
} // end parallel
```



Combined

OpenMP Worksharing Scheduling

Clause Syntax: `parallel for schedule(schedule-type[, chunk-size])`

Schedule Type

Schedule (**static**, **chunk**)

- Threads receive chunks of iterations in thread order, round-robin. (Divided “equally” if no chunk size ... N / p)
- Good if every iteration contains same amount of work ... uniform workload.
- May help keep parts of an array in a particular processor's cache—good between **parallel for**'s.

Schedule (**dynamic**, **chunk**)

- Thread receives chunks as it (the thread) becomes available for more work ... a queue.
- Default chunk size may be 1
- Good for load-balancing non-uniform workloads.

OpenMP Worksharing Scheduling

Schedule (**guided**, **chunk**)

- Thread receives chunks as the thread becomes available for work.
- Chunk size decreases exponentially, until it reaches the chunk size specified (default is 1).
- Balances load and reduces number of requests for more work.
- *(I have never found this useful.)*

Schedule (**runtime**)

- Schedule is determined at run-time by the OMP_SCHEDULE environment value.
- Useful for experimentation

OpenMP Worksharing Scheduling

For example, loop with 100 iterations and 4 threads

- schedule (**static**)

Thread	0	1	2	3
Iteration	1-25	26-50	51-75	76-100

- schedule (**dynamic**, 15) (*one possible outcome*)

Thread	0	1	3	2	1	3	2
Iteration	1-15	16-30	31-45	46-60	61-75	76-90	90-100

- Dynamic scheduling has overhead since assigning the “next” chunk requires thread communication (access to shared data object)

OpenMP worksharing -- Sections

- **sections**

- Blocks of code are split among threads - task parallel style
- A thread might execute more than one block or no blocks
- Implied barrier

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        TASK1 (...);
        #pragma omp section
        TASK2 (...);
        #pragma omp section
        TASK3 (...);
    } // <- implicit barrier
} // <- another barrier at the end of the parallel region
```

OpenMP worksharing - Single

- **single** (or master)
 - Block of code is executed only by the 1st thread (single) or rank-0 (master)
 - Implied barrier at '}' after single ... no barrier after master!

```
1  int global_count = 0; // Defined globally.
2  ...
3  #pragma omp parallel shared( global_count )
4  {
5      foo1();
6      #pragma single
7      {
8          global_count++;
9          printf("%d\n", global_count);
10     } // <- barrier
11     foo2();
12 }
```

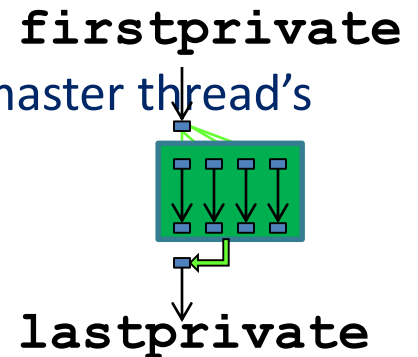

OpenMP Data scoping

`#pragma omp directive-name [clause [[,]clause]...]`

- **private** (variable list)
 - Each thread has its own copy of the specified variable
 - Variables are undefined after worksharing region
- **shared** (variable list)
 - Threads share a single copy of the specified variable
- **default** (type)
 - A default of **private**, **shared**, or **none** can be specified
 - Note that loop counter(s) of worksharing constructs are always **private** by default; everything else is **shared** by default

OpenMP Data Scoping

- **firstprivate** (variable list)
 - Like **private**, but copies are initialized with the master thread's value
- **lastprivate**(variable list)
 - Like **private**, but final value is copied out to master thread's copy
 - **for**: last iteration; **sections**: last section
- **reduction** (operator : variable)
 - Each thread has its own copy of the specified variable
 - Can appear only in reduction operation
 - All copies are "reduced" back into the original master thread's variable.
 - Example **reduction (+ : sum)** adds all values together.



OpenMP Data Scoping

- **for** and **parallel for** constructs
 - index variable is automatically private
- Automatic storage variables
 - private, if declared in a scope inside the construct (e.g. ordinary local variables declared inside functions)

OpenMP worksharing - Single

- **shared** - Variable is shared (seen) by all processors.
- **private** - Each thread has a private instance of the variable.
- Defaults: All loop indices are private, all other variables are shared.

```
double t1, t2;  
int i;  
#pragma parallel for shared(a) private(t1,t2)  
for (i = 0; i < 1000; i++){  
    t1 = f[i];  
    t2 = g[i];  
    a[i] = sqrt( t1*t1 + t2*t2 );  
}
```

OpenMP worksharing - Single

- **shared** - Variable is shared (seen) by all processors.
- **private** - Each thread has a private instance of the variable.
- Use local scoping whenever possible! It's a lot safer and unambiguous.

```
#pragma parallel for shared(a)
for (int i = 0; i < 1000; i++){
    double t1 = f[i];
    double t2 = g[i];
    a[i] = sqrt( t1*t1 + t2*t2 );
}
```

OpenMP Data Scoping

```
int sum = 0;
#pragma omp parallel for reduction( + : sum )
for (int i = 0; i < N; i++){
    sum = sum + a[i];
}
// Each thread's copy of sum is added
// to original sum at end of loop
printf("sum= %f\n",sum);

#pragma omp parallel for lastprivate( temp )
for (int i = 0; i < N; i++){
    temp = f[i];
}
printf("f[N-1] == %f\n", temp);
// temp is equal to f[N-1] at end of loop
```

OpenMP Synchronization

- **nowait** clause
 - Threads encounter a barrier synchronization at end of worksharing constructs.
 - Specifies that threads can proceed to the next task without waiting.

```
#pragma omp      ...   nowait
```

OpenMP Synchronization

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i=0; i<N; i++){
        b[i] = foo(a[i]);
    }
    // Threads can start foo2() before others finish foo()

    #pragma omp for
    for (i=0; i<N; i++){
        d[i] = foo2(c[i]);
    }
}
```

- Why? This avoids unnecessary synchronization ... which is overhead.
- Serial overhead increases the serial ratio in Amdahl's limit.

OpenMP Synchronization

- **barrier** directive explicitly adds synchronization point
 - All threads must reach the barrier and wait until the last gets there.

```
#pragma omp barrier
```

OpenMP Synchronization

```
#pragma omp parallel
{
    const int thread_id  = omp_get_thread_num(),
              num_threads = omp_get_num_threads();
    const int N_per_thread = N / num_threads;
    int istart = thread_id * N_per_thread,
        iend = istart + N_per_thread;

    for (i = istart; i < iend; i++)
        b[i] = foo( a[i] ); // Fixed cost per iteration

    // Don't let any threads past this point yet to avoid race
    // scenario on b[].
    #pragma omp barrier

    #pragma omp for schedule(dynamic) nowait
    for (int i = 0; i < N; i++){
        b[i] += foo2( c[i] );
    }
    // All threads block (or sync) here implicitly.
```

OpenMP Synchronization

- **critical** section permits access to only one thread at a time
 - All threads execute the critical section but they access the section with mutual exclusion ... only one can be inside the code at a given time.
 - This avoids race conditions but is very expensive.
- #pragma omp critical**

OpenMP Synchronization

```
InitializeQueue( allWorkItems ); // Initialize all work items serially.
#pragma omp parallel
{
    bool notEmpty = true; // This is a private variable for each thread.
    while ( notEmpty )
    {
        workItemType workItem;

        // Pop a task off the queue one thread at a time. This avoids
        // race condition on queue state.
        #pragma omp critical
        {
            notEmpty = PopQueue( &workItem );
        }

        // Do some work (if necessary)
        if ( notEmpty )
            foo(workItem);
    }
} // <= All threads block (or sync) here.
```

OpenMP Synchronization

- **atomic** statements insure that only one thread reads / writes / updates a shared variable memory location at a time.
 - These are generally faster than critical sections since they can be supported in hardware. And, if there isn't a conflict with another thread, this has little overhead.
 - Therefore, this is a cheap (scalable) strategy to avoid race conditions but is limited to simple statements.

```
#pragma omp atomic [read,write,capture]
```

OpenMP Synchronization

```
int nWorkItems = 4000, WorkItemCounter = 0;
WorkItemType workItem[] = InitializeItems(nWorkItems);
#pragma omp parallel shared( nWorkItems, workItemCounter )
{
    int myWorkItem = omp_get_thread_num();

    while (myWorkItem < nWorkItems)
    {
        foo( workItem[myWorkItem] );

        // copy the shared counter and then increment.
        #pragma omp atomic capture
        myWorkItem = WorkItemCounter++;
    }
} // <= All threads block (or sync) here.
```

- This is (usually) faster than a queue with a critical section since atomics are pretty quick.
- This is a way of manually coding a dynamic parallel for loop.