

COMP 364 / 464

High Performance Computing

Collection Commumcation:

Message Passing Interface (MPI)

MPI Collective Communications

- Every process *within the communicator* MUST call the routine
 - All calls are blocking
 - A process may return when participation is complete
 - May or may not synchronize (implementation dependent)
- You can build your own using p2p methods but MPI has done most of this for us (send/recv are the building blocks)
- Send and Receive sizes must match
 - mapping may vary
- Basic calls have a *root (or destination)* — “all” versions don’t

MPI Collective Communications

- Involves a group of processes.

- Basic Routines

Broadcast— `MPI_Bcast()`

Reduce— `MPI_Reduce()`

Gather/Scatter— `MPI_Gather()` `MPI_Scatter()` ...

- “All” versions

`MPI_Allreduce()`

`MPI_Allgather()`

`MPI_Alltoall()` ...

- Others

`MPI_Barrier` ...

Before

After

Root

“Send” array
element or
single variable

Task or
process

p0	A			
p1				
p2				
p3				

p0	A			
p1	B			
p2	C			
p3	D			

p0	A			
p1	B			
p2	C			
p3	D			

p0	A	B	C	D
p1				
p2				
p3				

Broadcast

Reduce

Gather

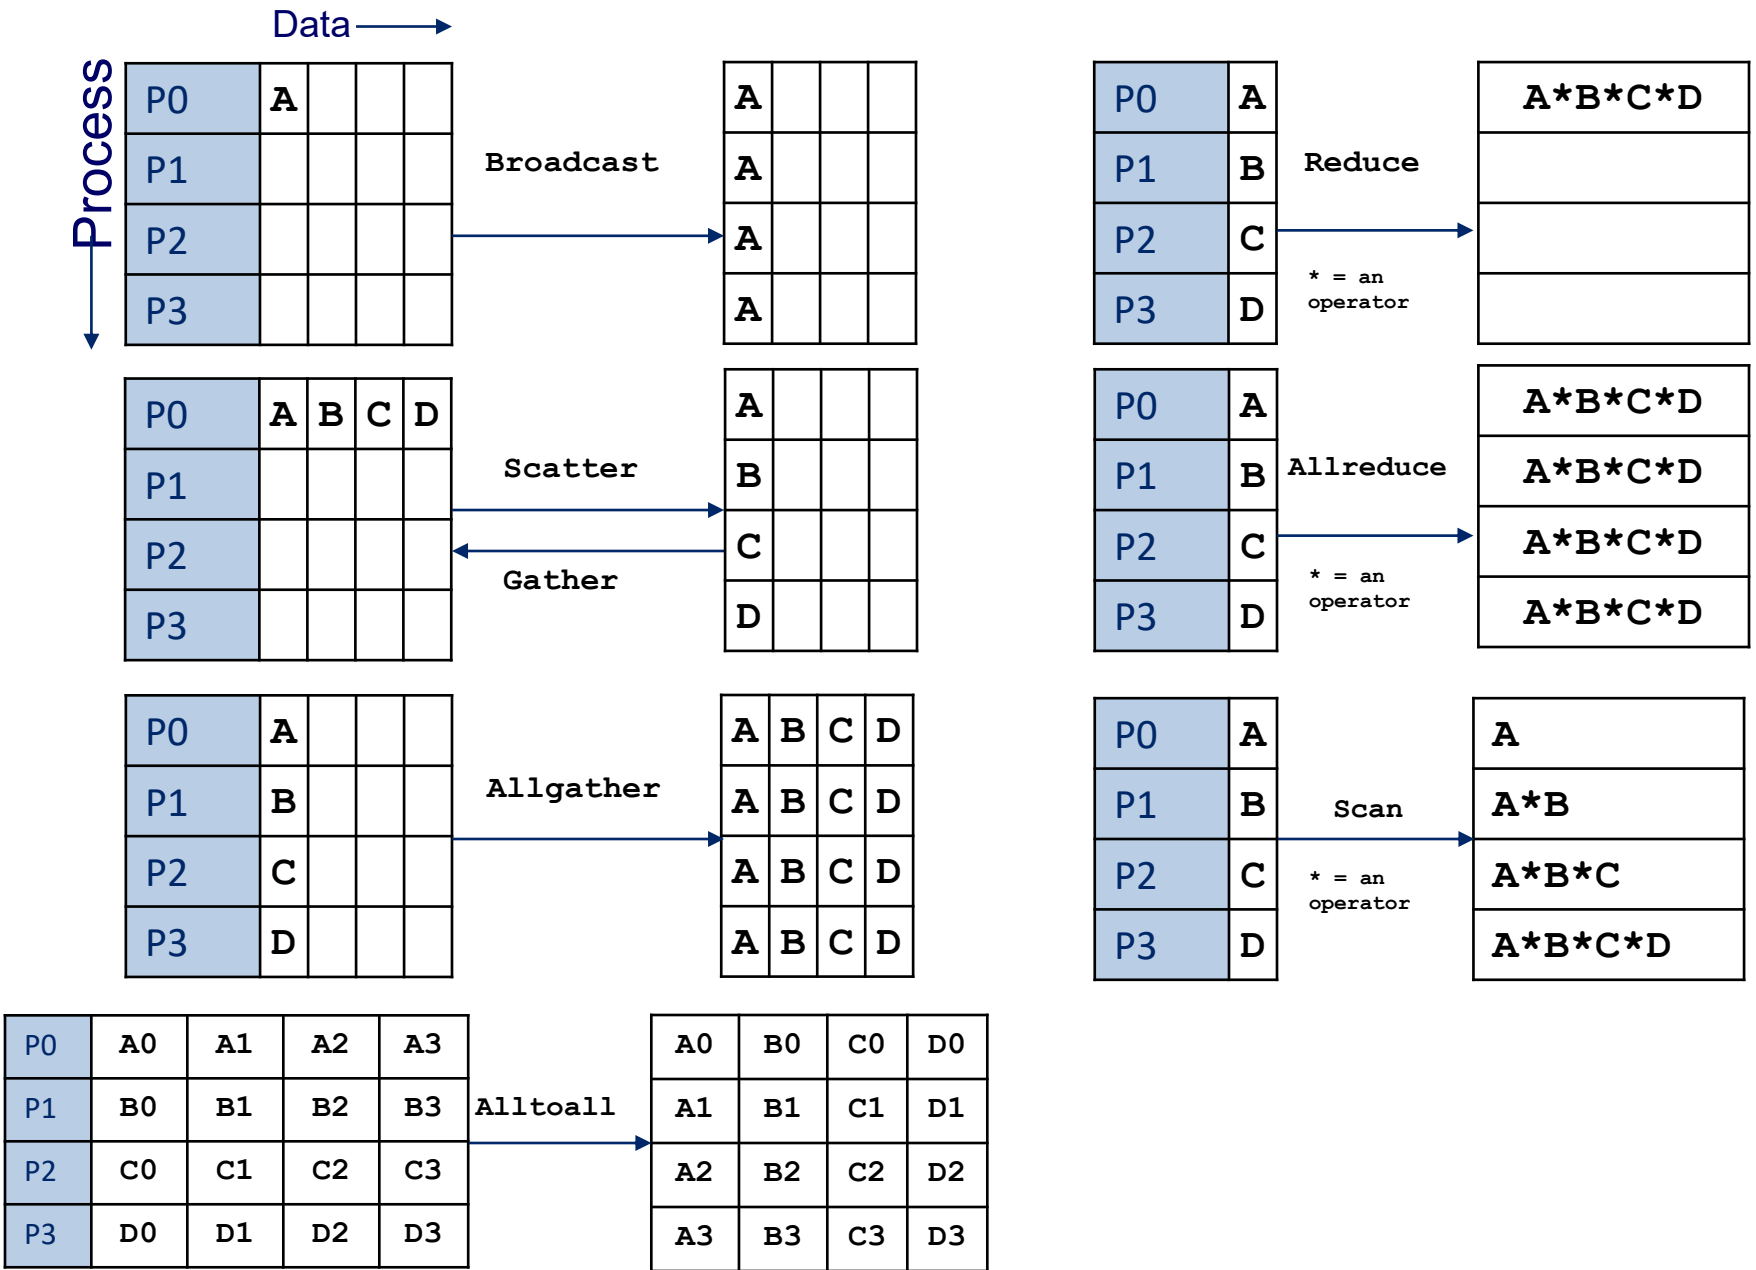
Scatter

p0	A			
p1	A			
p2	A			
p3	A			

p0	A op B op C op D			
p1				
p2				
p3				

p0	A	B	C	D
p1				
p2				
p3				

p0	A			
p1	B			
p2	C			
p3	D			



Broadcast Operation: `MPI_Bcast`

- All nodes call `MPI_Bcast`
- One node (`root`) sends a message to all
 - all others receive the message
- All MPI functions have error handling. `MPI_SUCCESS` = all good.

```
int ierr = MPI_Bcast(&dat, cnt, datatype, root,  
comm) ;
```

Reduction Operations

- Used to combine (reduce) partial results from all processors
- Result returned to root processor
- pre-defined or user-defined operations
 - Predefined: associative & commutative (com)
Order may not be canonical (i.e., rank order)
 - User defined: Must be associative. com or non-com
“Canonical” evaluation
- Works on a scalar variable or arrays (elemental)

MPI_Reduce

```
ierr = MPI_Reduce(&sbuf[0], &rbuf[0],  
count, datatype, operator, root, comm)
```

- Parameters
 - like **MPI_Bcast**, a root is specified
 - operation is a type of mathematical operation
- Applies the operator to each element globally
 - send and receive buffers are the same size
- Use **MPI_Op_create** for user-defined operation.

Operations for MPI_Reduce

MPI_PROD	Product
MPI_SUM	Sum
MPI_LAND	Logical and
MPI_LOR	Logical or
MPI_LXOR	Logical exclusive or
MPI_BAND	Bitwise and
MPI_BOR	Bitwise or
MPI_BXOR	Bitwise exclusive or
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_MAXLOC	Maximum value and location
MPI_MINLOC	Minimum value and location

Dot Product of Two Vectors

```
double a[N], b[N];  
  
...  
double localSum=0.0;  
for (int i = 0; i < N; ++i)  
    localSum += a[i]*b[i];  
  
double globalSum;  
const int root = 0;  
MPI_Reduce(&localSum, &globalSum, 1, MPI_DOUBLE,  
    MPI_SUM, root, MPI_COMM_WORLD);  
  
printf("rank= %d local= %f global= %f\n",  
    rank, localSum, globalSum);
```

MPI_Scatter Syntax

```
ierr = MPI_Scatter(&sbuf[0], scnt, stype, &rbuf[0],  
rcnt, rtype, root, comm);
```

- Parameters

- **sbuf** = array of size np*scnt (np = # of ranks)
- **scnt** = number of elements sent to each processor
- **rcnt** = number of element(s) obtained from the root processor
- **rbuf** = element(s) obtained from the root processor (rcnt in size)

e.g. MPI_Scatter(**S**, 1, stype, **R**, 1, rtype, root, comm)

Array

Scalar

We gratefully acknowledge the sponsorship of Chevron Corporation, whose generous support of TACC has made possible this Scientific Computing Curriculum and other student-focused initiatives.

© The University of Texas at Austin, 2013

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text:
"Parallel Computing for Science and Engineering course materials by The Texas Advanced Computing Center, 2013. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License"