

# COMP 364 / 464

## High Performance Computing

### **Distributed Memory Parallelism:** Point-to-Point (P2P) Communication

# Minimal MPI program

- Every MPI program needs these...
  - C version

```
#include <mpi.h>

...

ierr = MPI_Init(&argc, &argv);
ierr = MPI_Comm_size(MPI_COMM_WORLD, &numNodes);
ierr = MPI_Comm_rank(MPI_COMM_WORLD, &myRank);


...

ierr = MPI_Finalize();
```

- In C MPI routines are functions which return the error value

# Point-to-Point Communication

- Sending data from one point (process/task) to another point (process/task)
- One task **sends** while another **receives**



# Basic Communications in MPI

- Standard **MPI\_Send/MPI\_Recv** routines
  - Blocking calls used for basic messaging

## Point-to-Point Modes of Operation

- Blocking
  - Call does not return until the sent and received data is safe to use
- Non-blocking
  - Initiates send or receive operation, returns immediately
  - Can check or wait for completion of the operation
  - Data area is not safe for use until completion.
- Synchronous and Buffered (later)

# Data Types (basics)


- Data types (more of a mapping than a declaration)
  - Specifies the data type and size in MPI routines
  - Predefined MPI types correspond to language types

Representation	MPI Type	C
32-bit floating point	<code>MPI_FLOAT</code>	<code>float</code>
64-bit floating point	<code>MPI_DOUBLE</code>	<code>double</code>
32-bit integer	<code>MPI_INT</code>	<code>int</code>

- Methods exists for creating user-defined types
  - Simple (just combinations of normal data types)
  - Advanced (a map of data to be send)


# Life would be simple if....

- Processors would just send and receive, and the network would DWIM (do what I meant)




# Unfortunately...

- Data has to be somewhere: on one process or the other



# Non-Blocking Solution

- Create a buffer and let the send data sit there until someone picks it up






# Blocking Send/Receive

## Generic Syntax

- **MPI\_Send**(buf, count, datatype, dest, tag, comm)
- **MPI\_Recv**(buf, count, datatype, source, tag, comm, status)
- When MPI sends a message, it doesn't just send the contents; it also sends an *envelope* describing the contents:



Argument	Description
buf	initial address of send/receive buffer (reference): void*
count	number of items to send (integer): int
datatype	MPI data type of items to send/receive
dest	MPI rank of task receiving the data (integer): int
source	MPI rank of task sending the data (integer): int
tag	message ID (integer): int
comm	MPI communicator (set of exchange processors): MPI_Comm
status	returns information on the message received: MPI_Status

Parts of a P-2-P Communication:

**Data**  
**Send to/Recv from**  
**Message ID**

# Language Example

```
ierr = MPI_Send(&a[0], cnt, type, dest, tag, com) ;
```

```
ierr = MPI_Recv(&b[0], cnt, type, src, tag, com, &stat) ;
```

- Call blocks until send data of **a** has been sent *or copied to a buffer*.  
**recv**'s block until data is in **b**.

# P-2-P Example

```
#include <mpi.h>
int main(int argc, char* argv[])
{
    MPI_Comm Comm=MPI_COMM_WORLD;
    int numRanks,myRank=-1,ierr;

    ierr=MPI_Init(&argc, &argv);
    ierr=MPI_Comm_size(Comm,&numRanks);
    ierr=MPI_Comm_rank(Comm, &myRank);

    ierr=MPI_Finalize();

    printf("myRank=%d\n",myRank);
}
```

# P-2-P Example

```
#include <mpi.h>
int main(int argc, char* argv[])
{
    MPI_Comm Comm=MPI_COMM_WORLD; // Don't do this!
    MPI_Status status;
    int numRanks = 1, myRank = -1;
    int ierr=MPI_Init(&argc, &argv);
    ierr=MPI_Comm_size(Comm, &numRanks);
    ierr=MPI_Comm_rank(Comm, &myRank);

    int irecv = -1;
    if(myRank==0)
        ierr=MPI_Send(&myRank, 1,MPI_INT, 1,9, Comm);
    if(myRank==1)
        ierr=MPI_Recv(&irecv,1,MPI_INT, 0,9, Comm,&status);
    ierr=MPI_Finalize();

    printf("myRank=%d, received=%d\n",myRank,irecv);
}
```

# The 6 Basic MPI Call Summary

- MPI is used to create parallel programs based on message passing
- Usually the same program is run on multiple processors
- The 6 basic calls in MPI are:

```
MPI_Init(&argc, &argv) ;  
MPI_Comm_Rank (Comm, &myid) ;  
MPI_Comm_Size (Comm, &numprocs) ;  
MPI_Send(s_pointer, count, MPI_TYPE, dest, tag, Comm) ;  
MPI_Recv(r_pointer, count, MPI_TYPE, src, tag, Comm, &stat) ;  
MPI_Finalize() ;
```

**MPI\_TYPE** is an MPI Parameter

# MPI\_SendRecv

`MPI_SendRecv (sdat, scount, stype, dest, stag, rdat, rcount, rtype, src, rtag, comm, &status)`

- Initiates send and receive at the same time.
- Completes when both send and receive buffers are safe to use
- Useful for communications patterns where each node sends and receives messages (two-way communication). *Good for avoiding deadlock, implementing shifts/rings.*
- Executes a **standard mode** send & receive operation for **dest** and **src**, respectively.
- The send and receive operations use the same communicator, but have distinct tags.

# Blocking vs. Non-Blocking

- Blocking

- A blocking send routine will only return after it is *safe* to modify the data area.
- *Safe* means that modifications in the data area will not affect the data to be sent.
- A *Safe send* does not imply that the data was actually received.
- A blocking send can be either synchronous or asynchronous.

- Non-blocking

- Send/receive routines return immediately.
- Non-blocking operations request the MPI library to perform the operation when possible.
- It is **unsafe** to modify the data area until the requested operation has been performed. There are *wait* routines used to do this (**`MPI_Wait`**)
- Primarily used to overlap computation with communication

# Blocking vs. Non-Blocking Routines

Description	Syntax for C bindings
Blocking send	<code>MPI_Send(buf, count, datatype, dest, tag, comm)</code>
Non-blocking send	<code>MPI_Isend(buf, count, datatype, dest, tag, comm, request)</code>
Blocking receive	<code>MPI_Recv(buf, count, datatype, source, tag, comm, status)</code>
Non-blocking receive	<code>MPI_Irecv(buf, count, datatype, source, tag, comm, request)</code>
Wait for completion	<code>MPI_Wait(request, status)</code>

**request:** used by non-blocking send and receive operation



# Non-Blocking Communication

- Non-blocking send
  - send call returns immediately
  - send actually occurs later, call `mpi_wait()` to before destroying / modifying send data.
- Non-blocking receive
  - receive call returns immediately
  - when received data is needed, call a `wait()` to verify receipt
- Non-blocking communication used to overlap communication with computation (and communication with communication!).
- And ... used to prevent deadlock.

# Non-Blocking Send With **MPI\_Isend**

```
MPI_Request request;  
ierr = MPI_Isend(&data, count, datatype, dest, tag,  
                comm, &request);
```

- **request** is the id for the message call
- Don't use **data** area until communication is complete

# Non-Blocking Send With **MPI\_Irecv**

```
MPI_Request request;  
ierr = MPI_Irecv(&data, count, datatype, source,  
                tag, comm, &request);
```

- **request** is an id for communication
- Note: There is **no status parameter**
- Don't use **data** area until communication is complete

# MPI\_Wait Used to Complete Communication

- **request** from **MPI\_Isend** or **MPI\_Irecv**
  - the completion of a send operation indicates that the sender is now free to update the data in the send buffer
  - the completion of a receive operation indicates that the receive buffer contains the received message
- **MPI\_Wait** blocks until message specified by **request** completes

# MPI\_Wait Usage

```
MPI_Request request;  
MPI_Status status;  
...  
ierr = MPI_Wait(&request, &status)
```

# Non-Blocking Examples

# Two-way Communication: Deadlock

- Deadlock 1 (always deadlocks)

```
int other = 1 - myRank;    //!< Assume a total of 2 MPI tasks
MPI_Recv(recvbuf, count, MPI_FLOAT, other, tag, MPI_COMM_WORLD, status);
MPI_Send(sendbuf, count, MPI_FLOAT, other, tag, MPI_COMM_WORLD);
```

- Deadlock 2 (deadlocks when system buffer is too small)

```
other = 1-mytid    // Assume a total of 2 MPI tasks
MPI_Send(sendbuf, count, MPI_FLOAT, other, tag, MPI_COMM_WORLD);
MPI_Recv(recvbuf, count, MPI_FLOAT, other, tag, MPI_COMM_WORLD, status);
```

# Two-way Communication: Solutions

- Solution 1: specify sends & receives in an order that is guaranteed not to deadlock. This is cumbersome for more than two processors

```
if (rank==0) then
    MPI_Send(sendbuf,count,MPI_FLOAT,1,tag,MPI_COMM_WORLD)
    MPI_Recv(recvbuf,count,MPI_FLOAT,1,tag,MPI_COMM_WORLD,status)
elseif (rank==1) then
    MPI_Recv(recvbuf,count,MPI_FLOAT,0,tag,MPI_COMM_WORLD,status)
    MPI_Send(sendbuf,count,MPI_FLOAT,0,tag,MPI_COMM_WORLD)
endif
```

- Solution 2: use a sendrecv instruction. This works for communication where every processors does one send and one receive.

```
other = 1-myRank;      //!< Assume there are exactly 2 total tasks
MPI_Sendrecv(    sendbuf,sendcount,sendtype,other,sendtag,
    recvbuf,recvcount, recvtype,other,recvtag,MPI_COMM_WORLD,&status)
```



# Two-way Communication: Solutions

- Solution 3: use `isend` and `irecv`.
  - This is easier to write and probably more efficient than solution 1.
  - It can deal with more general communication patterns than solution 2.

```
MPI_Isend(sendbuf, count, MPI_FLOAT, other, tag, MPI_COMM_WORLD, &req1);  
MPI_Irecv(recvbuf, count, MPI_FLOAT, other, tag, MPI_COMM_WORLD, &req2);  
MPI_Wait(req1, &status);  
MPI_Wait(req2, &status);
```

# Two-way Communications Summary

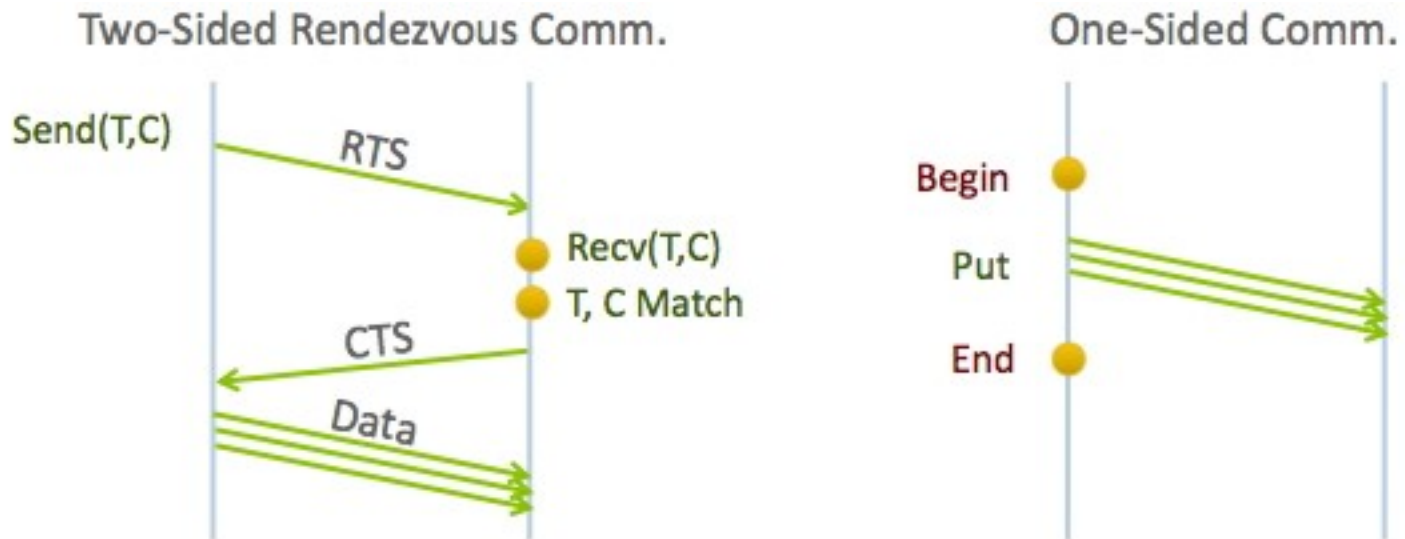
	CPU 1	CPU 2
Deadlock 1	Recv / Send	Recv / Send
Deadlock 2	Send / Recv	Send / Recv
Solution 1	Send / Recv	Recv / Send
Solution 2	SendRecv	SendRecv
Solution 3	Isend / Irecv/ Wait	Isend/ Irecv/ Wait

# Wait Types

- **MPI\_Wait** : wait for one request
- **MPI\_Waitall** : wait for an array of requests, good for load balanced tasks, or when all needed
- **MPI\_Waitany** : wait for one in an array of requests, good for unbalanced tasks, or if they can be processed individually
- **MPI\_Waitsome** : wait for any number in an array, much like **MPI\_Waitany**

# One-Sided Communications

- It would be nice to avoid that two-way orchestration: just write into another process' memory or read from it
- Less overhead, easier to code



# Wildcards (C)

- Enables programmer to avoid having to specify a tag and/or source.
- Example:

```
MPI_Status status;  
int data[5];  
int ierr;  
ierr = MPI_Recv(&data[0], 5, MPI_INT,  
               MPI_ANY_SOURCE, MPI_ANY_TAG,  
               MPI_COMM_WORLD, &status);
```

- `MPI_ANY_SOURCE` and `MPI_ANY_TAG` are wild cards
- `status` structure is used to get wildcard values

# More on Status

- **status** (type **MPI\_Status**) is a structure which contains three fields **MPI\_SOURCE**, **MPI\_TAG**, and **MPI\_ERROR**
- **status.MPI\_SOURCE**, **status.MPI\_TAG**, and **status.MPI\_ERROR** contain the source, tag, and error code respectively of the received message

# Order Semantics

- Messages with the same tag are ordered; for the rest, make no assumptions on message ordering!

- the first receive always matches the first send in the following

```
tag=123456
```

```
if (rank == 0) then
```

```
    call MPI_Send(b1,cnt,MPI_REAL,1,tag,comm,err)
```

```
    call MPI_Send(b2,cnt,MPI_REAL,1,tag,comm,err)
```

```
ELSE ! rank.EQ.1
```

```
    call MPI_Recv(b1,cnt,MPI_REAL,0,tag,comm,  
                  status,ierr)
```

```
    call MPI_Recv(b2,cnt,MPI_REAL,0,tag,comm,  
                  status, ierr)
```

```
END if
```

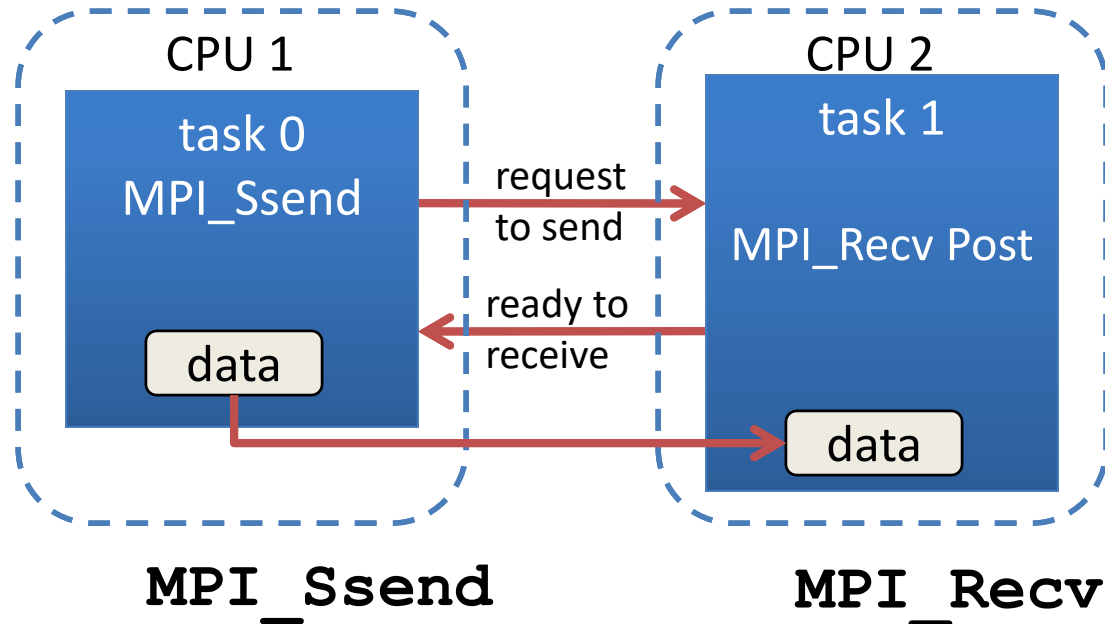
# MPI Receive Modes

**IMPORTANT: From the MPI-2.2 Standard:**

"There is only one receive operation, but it matches any of the send modes [emphasis mine]. The receive operation described in the last section is blocking: it returns only after the receive buffer contains the newly received message. A receive can complete before the matching send has completed (of course, it can complete only after the matching send has started)."



# Synchronous Communication



- Data isn't sent until Receive has been posted.
- Synchronous send returns when data area is safe for re-use.
- There is no **MPI\_Srecv**

# Synchronous Communication

## Ssend

```
...  
i=1;  
if(irank == 0){  
    MPI_Ssend(&i, 1, MPI_INT,      1, 9, MPI_COMM_WORLD);  
}else {  
    MPI_Recv( &j, 1, MPI_INT,      0, 9, MPI_COMM_WORLD, &status);  
}
```

# MPI\_Test

- Value of flags signifies whether a message has been delivered
- Similar to **MPI\_Wait**, but does not block

```
int flag;
```

```
ierr= MPI_Test(&request, &flag, &status);
```

# MPI\_Cancel

- Cancel a pending non-blocking send or receive

```
MPI_Request request;  
ierr= MPI_Cancel(&request);
```

# MPI\_Probe

- **MPI\_Probe** allows incoming messages to be checked without actually receiving them
  - the user can then decide how to receive the data
  - Used when different actions need to be taken, depending on the "who, what, and how much" information of the message.

# MPI\_Probe

```
ierr = MPI_Probe(source, tag, comm, &status);
```

- Parameters
  - source: source rank or **MPI\_ANY\_SOURCE**
  - tag: tag value or **MPI\_ANY\_TAG**
  - comm: communicator
  - status: status object

*We gratefully acknowledge the sponsorship of Chevron Corporation, whose generous support of TACC has made possible this Scientific Computing Curriculum and other student-focused initiatives.*

© The University of Texas at Austin, 2013

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text:  
*"Parallel Computing for Science and Engineering course materials by The Texas Advanced Computing Center, 2013. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License"*