

1. A system has an 8-way set associative 32KB L1 data cache with 64B cachelines.

a. How many cachelines are in the L1 cache?

b. How many C/C++ double words can fit in each cacheline?

- c. Examine the following code. Suggest a modification that could improve the cache efficiency on this system and explain why.

```
double x[16][512], y[512];

for (int i = 0; i < 512; ++i)
    for (int j = 0; j < 16; ++j)
        y[i] = y[i] + x[j][i];
```

2. Examine the two matrix multiplication methods below and suggest which one you would expect to be faster on a CPU with a cache. (This may sound redundant but some CPUs didn't have caches historically or even in modern times.) Assume the input matrices (A,B) are square with N rows and columns and N is a multiple of 4. Explain your answer.

a.

```
double A[N][N], B[N][N], C[N][N];
for (int i = 0; i < N; ++i)
    for (int j = 0; j < N; ++j) {
        C[i][j] = 0.0;
        for (int k = 0; k < N; ++k)
            C[i][j] = C[i][j] + A[i][k] * B[k][j];
    }
```

b.

```
for (int i = 0; i < N; ++i) {
    for (int j = 0; j < N; ++j) C[i][j] = 0.0;

    for (int k = 0; k < N; k = k+2) {
        double a0 = A[i][k];
        double a1 = A[i][k+1];

        for (int j = 0; j < N; ++j) {
            double cij = C[i][j];
            cij = cij + a0 * B[k][j];
            cij = cij + a1 * B[k+1][j];
            C[i][j] = cij;
        }
    }
}
```